

# DARSHAN ASHOK FOFAJIYA

Senior Applied Scientist – Amazon

253-368-9607 | [fofadiyadarshan@gmail.com](mailto:fofadiyadarshan@gmail.com) | [linkedin.com/in/darshan-fofadiya](https://linkedin.com/in/darshan-fofadiya) | Seattle, WA

## RESEARCH SUMMARY

Senior Applied Scientist with **11+ years of experience** architecting large-scale ML systems. Specializing in **High-Performance Computing (HPC)**, **GPU optimization**, and **non-convex optimization landscapes**. Bridging the gap between theoretical guarantees and empirical scale by designing custom optimizers (GatedAdam) and novel architectures (Idea-Gated Transformers). Combining industrial rigor with independent research to solve stability and trainability in deep networks.

## INDEPENDENT RESEARCH & NOVEL ARCHITECTURES

### Idea-Gated Transformers: Latent Semantic Planning | Under Review, ICML 2026

2025 – Present

- **Novel Architecture:** Proposed a Transformer variant that separates "concept planning" from "token generation" using differentiable soft-gating. This mitigates **associative drift** in long-horizon generation.
- **Relevance to Physics:** The separation of "semantic state" from "syntactic generation" mirrors the need to maintain quantum state coherence against decoherence drift.
- **Implementation:** Built custom PyTorch kernels to handle the gating mechanism efficiently on GPU.

### GatedAdam: Biologically Inspired Sparse Optimization | Preprint

2025 – Present

- **Solving Gradient Vanishing:** Developed a custom optimizer to tackle the **vanishing gradient problem** (analogous to Barren Plateaus) by using activation saturation as a constructive signal.
- **Mechanism:** Implemented a "Benefit Score" gating mechanism ( $O(1)$  overhead) to enforce **93.43% structural sparsity**, replacing standard regularization with a physical "braking system."
- **Stability:** Stabilized training in non-convex regimes where standard Adam diverged.

## PROFESSIONAL EXPERIENCE

### Amazon

Seattle, WA

#### Senior Applied Scientist (L6)

Oct 2023 – Present

- **HPC & Distributed Systems:** Architected a custom density-based clustering engine for **450k high-dimensional entities** across 20+ languages. Solved **topological collapse** in embedding spaces by engineering an "exponential growth trigger" with binary search, stabilizing graph construction where standard DBSCAN failed.
- **GPU Optimization:** Optimized GPU memory utilization for distributed training runs, reducing compute cost by 40% while increasing batch size throughput.
- **Causal Inference:** Applied Double Machine Learning (DML) to assess causal uplift, influencing strategic resource allocation for Amazon Books.

### Amazon

Seattle, WA

#### Applied Scientist II

June 2018 – Sep 2023

- **Robustness & Noise Tolerance:** Owned the architecture for **False Reject Detection** in Alexa. Developed semi-supervised frameworks to robustly learn from noisy, high-variance data in low-resource domains.
- **Scale:** Deployed production-grade NER and Intent Classification models serving millions of daily requests with strict latency constraints.
- Published research in **Findings of the ACL (2021)**.

### ZS Associates

Pune, India / Evanston, IL

#### Data Scientist

May 2015 – May 2018

- Built **embedding representations** for all US physicians to improve targeting and segmentation, preceding the widespread adoption of transformer-based embeddings in the pharma domain.

### Mu Sigma

Bangalore, India

#### Decision Scientist

Aug 2014 – May 2015

- Worked on customer segmentation and targeting problems for major US retail and pharma clients.

## PUBLICATIONS & PATENTS

**Preprint:** Idea Gated Transformers: Latent Semantic Planning. arXiv:2512.03343 (Submitted ICML 2026).

**Paper:** False Reject Detection for Low-Resource Domains in Neural Spoken Language Understanding. Findings of the ACL-IJCNLP 2021.

**Patent US11604925:** Architecture for gazetteer-augmented named entity recognition. Issued Mar 2023.

**Patent US11823671:** Architecture for context-augmented word embedding. Issued Nov 2023.

## EDUCATION

---

### University of Mumbai

*Bachelor of Engineering in Electrical Engineering*

Mumbai, India

Aug 2010 – May 2014

- Focus on Complex Analysis, Advanced Mathematics, and Control Systems.

## TECHNICAL SKILLS

---

**HPC & Frameworks:** PyTorch (Advanced), CUDA Optimization, Distributed Data Parallel (DDP), JAX

**Optimization:** Stochastic Gradient Descent Dynamics, Loss Landscape Analysis, Custom Optimizer Implementation

**Algorithms:** Graph Theory, Transformer Architectures, Reinforcement Learning (PPO), Density Clustering

**Languages:** Python, C++, SQL